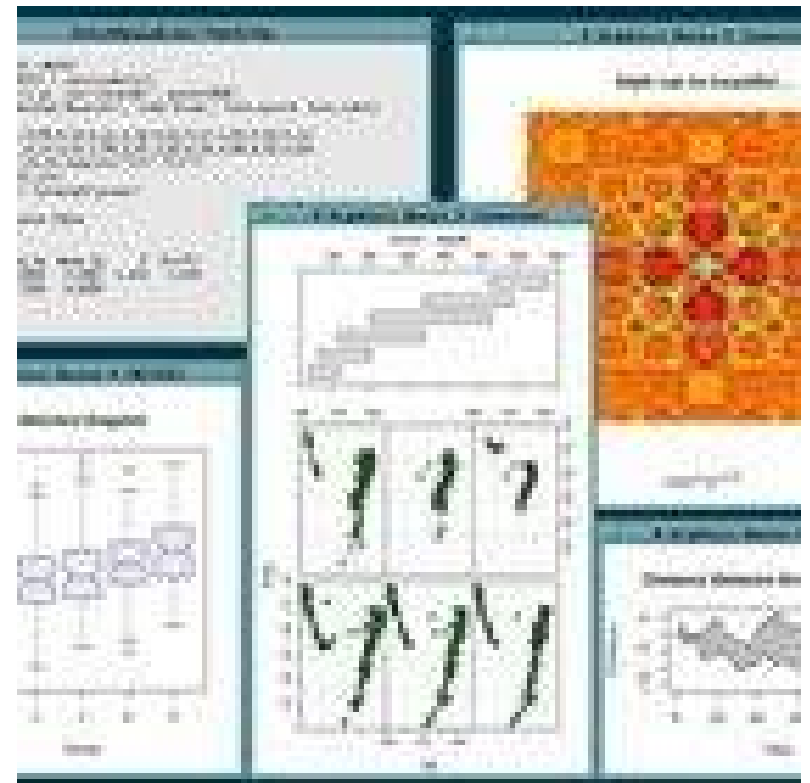
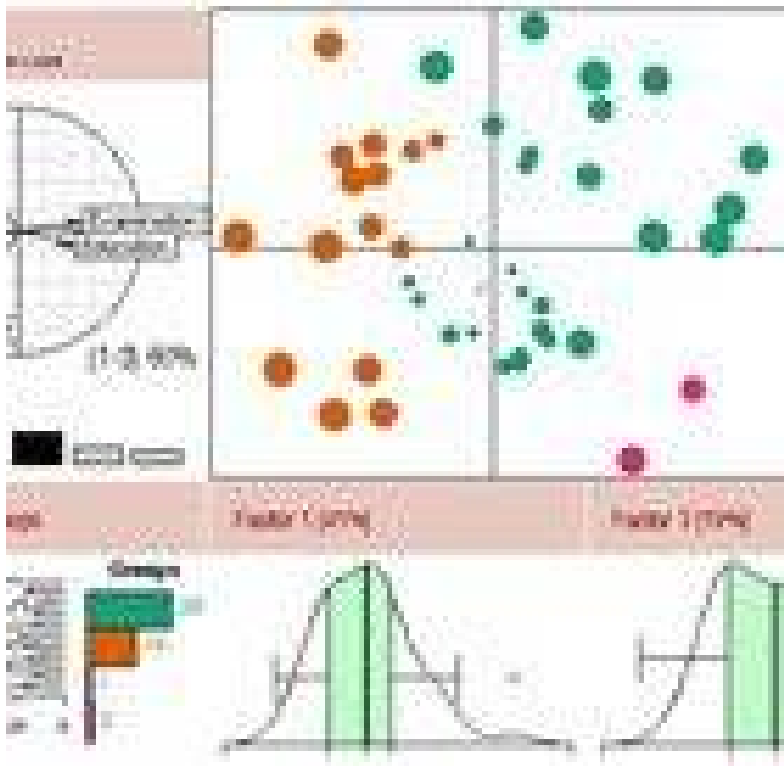


SCS2111 – Laboratory II

Lecture 01

19.08.2015



SCS2111 – Laboratory II

- 1L + 1P (Compulsary)
- Lecturers :
 - Dr. C. H. Magalla (champa@stat.cmb.ac.lk)
 - Mrs. Mindika Premachandra (amp@ucsc.cmb.ac.lk)
- Lectures : 1.00 – 2.00 pm Thursdays (Mini Aud.)
- Practicals :
 - 3.00 – 5.00 pm Wednesdays (IRQUE)
 - 2.00 – 5.00 pm Thursdays (Labs A & B, IRQUE)

Learning Outcomes

- Be familiar with Matlab/Octave and R environments
 - Be able to do basic mathematics and simple manipulations with Matlab and R
 - Be able to do basic statistical operations with R
 - Use Matlab and R for basic plotting
 - Learn scripting with Matlab and R
 - Use different data analysis techniques in Matlab and R

Course Contents : Matlab/Octave

- Manipulating Variables
- Basic Mathematics using Matlab
 - E.g.: Linear Algebra, basic statistics, differentiation and integrals, Fourier transforms
- Basic plotting and curve fitting
- Programming Scripts and Functions
- Image Processing functions and Animations
- Debugging
- Data Structures and File Management
- Symbolics, Simulink, file I/O, building GUIs

Course Contents: R

- Introduction to R environment, Getting Help, R Commands, Case sensitivity, Recall and correction of previous commands, R scripting and executing
- Simple manipulations: objects, vector arithmetic, arrays and matrices, lists and data frames
- Reading data from file
- R as a statistical package: Analysis of Qualitative and Quantitative data, probability distributions, Numerical measures, probability distributions

Course Contents : R

- Graphical procedures: plotting, displaying multivariate data, display graphics
- Multivariate correlations: Bayes approach and Naïve Bayes classifier
- Principal Component Analysis
- Cluster Analysis: K means clustering, K nearest neighbours clustering
- Linear Discriminant Analysis
- Decision Tree Models (Tree-based) models

Assessment

- Lab sheets
- Assignments
- Final Examination
- Rubric : 70% Theory, 30% Practicals?

Detailed assessment plan will be informed later

What is R?

- R is a free software environment for statistical computing and graphics
- R was created by [Ross Ihaka](#) and [Robert Gentleman](#) at the University of Auckland
- A GNU project which is similar to the **S** language,
 - developed at Bell Laboratories by [John Chambers](#) and colleagues.
 - **R** can be considered as a different implementation of **S**.
- The source code for the R software environment is written primarily in [C](#), [Fortran](#), and R.
- Pre-compiled binary versions are provided for various OS
- <http://www.r-project.org/index.html>

Why learn R?

- R is FREE, easy to use, and open source.
 - Commercial options: SAS, SPSS
- The R language is widely used among statisticians and data miners for developing statistical software and data analysis
- The "de facto" standard for data analysis and data mining
- A complete programming language
- Comes with a large library of pre-defined functions
- Better suited for advanced users who want all the power in their hands
 - R supports [matrix arithmetic](#)
 - R's [data structures](#) include [vectors](#), [matrices](#), arrays, data frames (similar to [tables](#) in a [relational database](#)) and [lists](#).
 - R's extensible object system includes objects for (among others): [regression models](#), [time-series](#) and [geo-spatial coordinates](#).



The 2015 Top Ten **Programming Languages**

IEEE Spectrum - Jul 20, 2015

What are the most popular **programming languages**? ... The big mover is **R**, a statistical computing language that's handy for analyzing and ...



R Rises in IEEE Ranking of Top **Programming Languages**

ADT Magazine - Jul 21, 2015


IEEE Spectrum has followed up last year's report on the top **programming languages** with a new study that sees **R** making a big jump in the ...



In data science, **the R language** is swallowing Python

InfoWorld - Jul 24, 2015

It's always precarious to compare **programming languages**, given their ... While **R** is a language developed by and for statisticians, Python has ...




The Most Popular **Programming Languages** of 2015

ProgrammableWeb - Aug 4, 2015

While the top 5 remain unchanged, C has moved within touching distance of Java, and statistical **programming language R** has jumped from ...

Learning Resources

- An excellent introductory book is by Peter Dalgaard, "Introductory Statistics with R", Springer (2002)
- *An Introduction to R. Notes on R: A Programming Environment for Data Analysis and Graphics. Version 2.15.2 (2012-10-26). W. N. Venables, D. M. Smith.*



Interacting with R

- R is an interpreted language; users typically access it through a command-line interpreter.
- There are also several graphical front-ends for it.
- Unlike languages like C, Fortran, or Java, R is an interactive programming language.
- This means that R works interactively, using a question-and-answer model:
 - Start **R**
 - Type a command and press **Enter**
 - **R** executes this command (often printing the result)
 - **R** then waits for more input
 - Type **q()** to exit

Here are some simple examples:

- Taken from AN INTRODUCTION TO R by Deepayan Sarkar

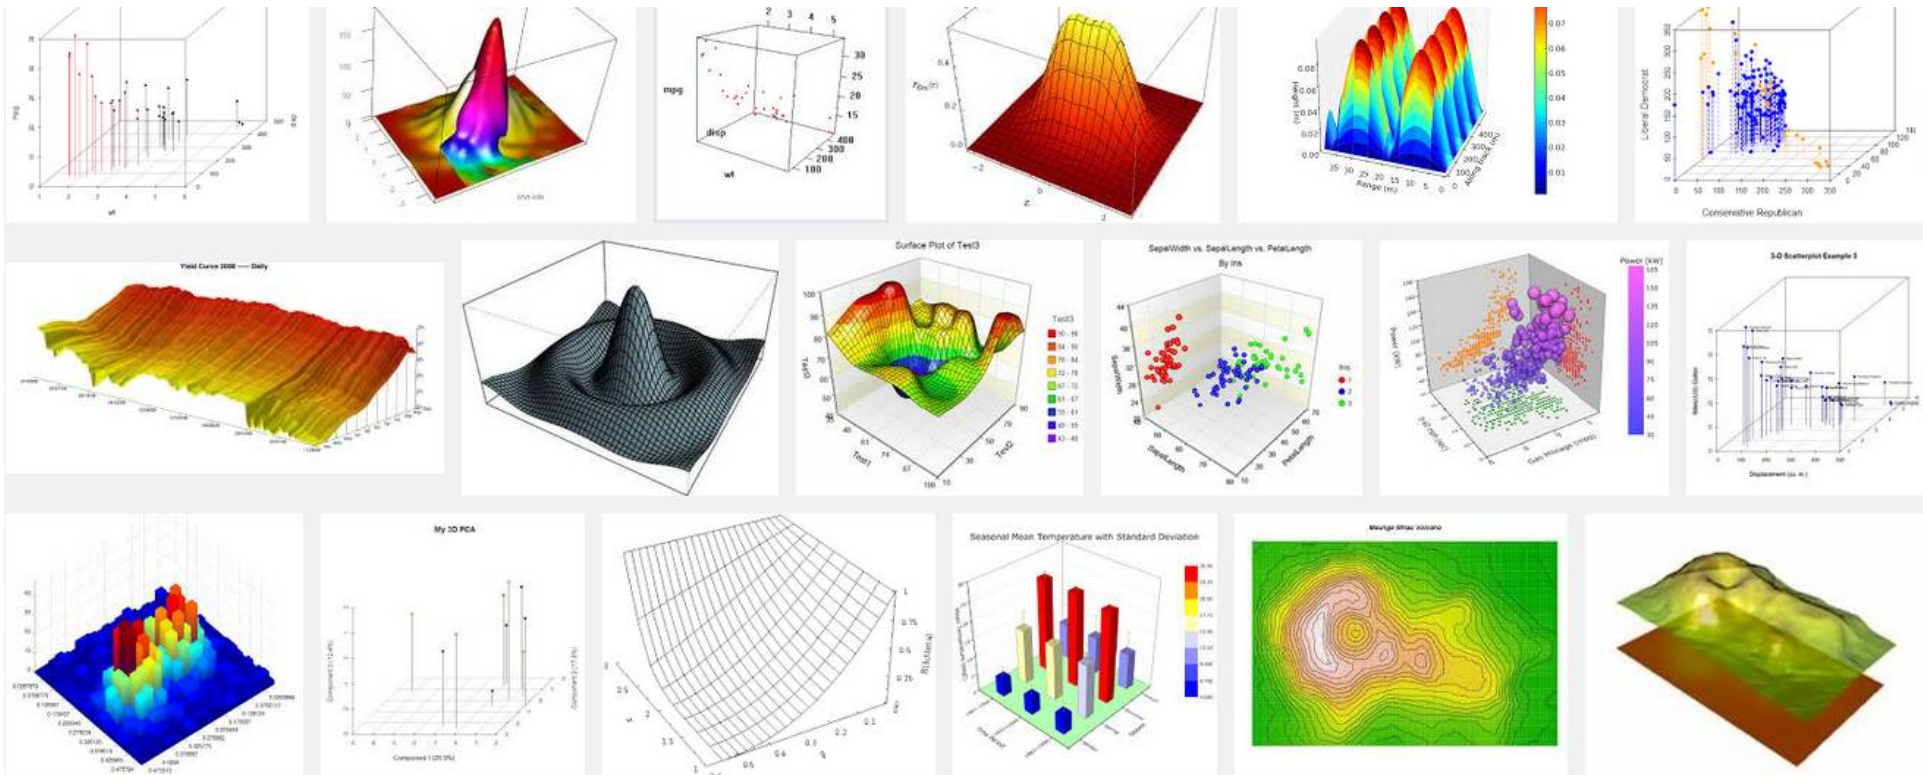

```
> 2 + 2
[1] 4
> exp(-2) ## exponential function
[1] 0.1353353
> log(100, base = 10)

[1] 2
> runif(10)
[1] 0.39435130 0.98811744 0.07357143 0.16689946 0.80572031 0.05292909
[7] 0.70498250 0.18781961 0.07865185 0.21618324
```

- The last command generates ten $U(0; 1)$ random variables; the result (which is printed) is a vector of 10 numbers. `exp()`, `log()`, and `runif()` are functions.

Plots

```
> plot(runif(10))
```



Variables

- R has symbolic variables which can be assigned values.
- Assignment is done using the '<-' operator.
- The more C-like '=' also works (with some exceptions).

```
> s <- "this is a character string"
> s
[1] "this is a character string"
```

```
> x <- 2
> x + x
[1] 4
> yVar2 = x + 3
> yVar2
[1] 5
```

- Variable names can be almost anything, but they should not start with a digit, and should not contain spaces. Names are case-sensitive.
- Some common names are already used by R (c, q, t, C, D, F, l, T) and should be avoided.

Vectorized arithmetic

- The elementary data types in R are all vectors; even the scalar variables we defined above are stored as vectors of length one.
- The **c(...)** construct can be used to create vectors:

```
> weight <- c(60, 72, 57, 90, 95, 72)
> weight
[1] 60 72 57 90 95 72
```

- To generate a vector of regularly spaced numbers, use

```
> seq(0, 1, length = 11)
[1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
> 1:10
[1] 1 2 3 4 5 6 7 8 9 10
```


- The `c()` function can be used to combine vectors as well as scalars,

```
> x <- seq(0, 1, length = 6)
> c(x, 1:10, 100)

[1] 0.0 0.2 0.4 0.6 0.8 1.0 1.0 2.0 3.0 4.0 5.0 6.0
[13] 7.0 8.0 9.0 10.0 100.0
```

- Common arithmetic operations

(including `+`, `-`, `*`, `/`, `^`) and mathematical functions (e.g. `sin()`, `cos()`, `log()`) work element-wise on vectors, producing another vector:

```
> height <- c(1.75, 1.80, 1.65, 1.90, 1.74, 1.91)
> height^2
[1] 3.0625 3.2400 2.7225 3.6100 3.0276 3.6481
> bmi <- weight / height^2
> bmi
[1] 19.59184 22.22222 20.93664 24.93075 31.37799 19.736
> log(bmi)
[1] 2.975113 3.101093 3.041501 3.216102 3.446107 2.9824
```

- When two vectors are not of equal length, the shorter one is recycled.
 - E.g.: The following adds 0 to all the odd elements and 2 to all the even elements of `1:10`:

```
> 1:10 + c(0, 2)

[1] 1 4 3 6 5 8 7 10 9 12
```

Summaries

- Many functions summarize a data vector by producing a scalar from a vector, e.g.,
- Simple summary statistics (mean, median, s.d., variance) can be computed from numeric vectors using appropriately named functions:

```
> sum(weight)
[1] 446
> length(weight)
[1] 6
> avg.weight <- mean(weight)
> avg.weight
[1] 74.33333
```

```
> x <- rnorm(100)
> mean(x)
[1] -0.1354077
> sd(x)
[1] 1.007307
> var(x)
[1] 1.014668
> median(x)
[1] -0.06083453
```

- Quantiles can be computed using the `quantile()` function.

- `IQR()` computes the inter-quartile range (**midspread** or **middle fifty**).

```
> xquants <- quantile(x)
> xquants
```

	0%	25%	50%	75%	100%
	-3.14440776	-0.74831291	-0.06083453	0.50980136	2.19369423

```
> xquants[4] - xquants[2]
```

	75%
	1.258114

```
> IQR(x)
```

```
[1] 1.258114
```

```
> quantile(x, probs = c(0.2, 0.4, 0.6, 0.8))
```

	20%	40%	60%	80%
	-1.0308886	-0.4388473	0.1236059	0.7357803

- The five-number summary (minimum, maximum, and quartiles) is given by `fivenum()`. A slightly extended summary is given by `summary()`.

```
> fivenum(x)
```

```
[1] -3.14440776 -0.75013378 -0.06083453 0.51742360 2.19369423
```

```
> summary(x)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-3.14400	-0.74830	-0.06083	-0.13540	0.50980	2.19400

Object-oriented programming: classes and methods

- Let's illustrate using a real dataset, one of the many datasets built into R (The well-known [Iris data](#)).
 - The dataset contains measurements on 150 flowers, 50 each from 3 species: Iris setosa, versicolor and virginica.
 - It is typically used to illustrate the problem of classification : [Given the four measurements for a new flower, can we predict its Species?](#)
 - Like most datasets, iris is not a simple vector, but a composite ["data frame"](#) object made up of several component vectors.
 - We can think of a data frame as a matrix-like object, with each row representing an observational unit (in this case, a flower), and columns representing multiple measurements made on the unit.

- The Iris data : The `head()` function extracts the first few rows, and the `$` operator extracts individual components.

```
> head(iris) # The first few rows
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
> iris$Sepal.Length
```

```
[1] 5.1 4.9 4.7 4.6 5.0 5.4 4.6 5.0 4.4 4.9 5.4 4.8 4.8 4.3 5.8 5.7 5.4 5.1
[19] 5.7 5.1 5.4 5.1 4.6 5.1 4.8 5.0 5.0 5.2 5.2 4.7 4.8 5.4 5.2 5.5 4.9 5.0
[37] 5.5 4.9 4.4 5.1 5.0 4.5 4.4 5.0 5.1 4.8 5.1 4.6 5.3 5.0 7.0 6.4 6.9 5.5
[55] 6.5 5.7 6.3 4.9 6.6 5.2 5.0 5.9 6.0 6.1 5.6 6.7 5.6 5.8 6.2 5.6 5.9 6.1
[73] 6.3 6.1 6.4 6.6 6.8 6.7 6.0 5.7 5.5 5.5 5.8 6.0 5.4 6.0 6.7 6.3 5.6 5.5
[91] 5.5 6.1 5.8 5.0 5.6 5.7 5.7 6.2 5.1 5.7 6.3 5.8 7.1 6.3 6.5 7.6 4.9 7.3
[109] 6.7 7.2 6.5 6.4 6.8 5.7 5.8 6.4 6.5 7.7 7.7 6.0 6.9 5.6 7.7 6.3 6.7 7.2
[127] 6.2 6.1 6.4 7.2 7.4 7.9 6.4 6.3 6.1 7.7 6.3 6.4 6.0 6.9 6.7 6.9 5.8 6.8
[145] 6.7 6.7 6.3 6.5 6.2 5.9
```

- A more concise description is given by the `str()` function (short for “structure”).

```
> str(iris)
'data.frame':      150 obs. of  5 variables:
 $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1
```

- As we can see,
 - the first four components of iris are numeric vectors,
 - but the last is a “`factor`”. These are how R represents categorical variables.

- Let us now see the effect of calling `summary()` for different types of objects.

- Note the different formats of the output.
- Species is summarized by the frequency distribution of its values because it is a categorical variable, for which mean or quantiles are meaningless.

- The entire data frame iris is summarized by combining the summaries of all its components.

```
> summary(iris$Sepal.Length)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
4.300  5.100   5.800   5.843  6.400   7.900

> summary(iris$Species)
   setosa versicolor  virginica 
      50         50       50 

> summary(iris)
  Sepal.Length  Sepal.Width  Petal.Length  Petal.Width
Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
Median :5.800   Median :3.000   Median :4.350   Median :1.300
Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500

  Species
setosa   :50
versicolor:50
virginica :50
```


- R achieves this kind of object-specific customized output through a fairly simple object-oriented paradigm.
- Each R object has a class ("numeric", "factor", etc.).

```
> class(iris$Sepal.Length)
[1] "numeric"
> class(iris$Species)
[1] "factor"
> class(iris)
[1] "data.frame"
```

- `summary()` is what is referred to as a generic function, with class-specific methods that handle objects of various classes.
- When the generic `summary()` is called, R figures out the appropriate method and calls it.

- The rules are fairly intuitive.
- The last call gives the list of all available methods.

```
> methods(summary)

[1] summary.aov                summary.aovlist             summary.aspell*
[4] summary.connection         summary.data.frame          summary.Date
[7] summary.default            summary.ecdf*               summary.factor
[10] summary.glm                summary.infl                summary.lm
[13] summary.loess*             summary.manova              summary.matrix
[16] summary.mlm                summary.nls*                summary.packageStatus*
[19] summary.PDF_Dictionary*    summary.PDF_Stream*        summary.POSIXct
[22] summary.POSIXlt            summary.ppr*                summary.prcomp*
[25] summary.princomp*          summary.srcfile             summary.scref
[28] summary.stepfun            summary.stl*                summary.table
[31] summary.tukeysmooth*

Non-visible functions are asterisked
```

- Objects of class “factor” are handled by [summary.factor\(\)](#),
- “data.frame”s are handled by [summary.data.frame\(\)](#).
- There is no [summary.numeric\(\)](#), so numeric vectors are handled by [summary.default\(\)](#).

Getting Help

- `help.start()` starts a browser window with an HTML help interface.
- `help(topic)` displays the help page for a particular topic. Every R function has a help page.

```
> help(plot)
> ?plot
```

```
> help(plot, help_type = "html")
```

- `help.search("search string")` performs a subject/keyword search.

```
> help.search("logarithm")
> ??logarithm
```

- To directly run the examples given in help pages, use the `example()` function

```
> example(plot)
```

- The `apropos()` function, lists all functions (or other variables) whose name matches a specified character string.

```
> apropos("plot")
[1] "assocplot"      "barplot"      "barplot.default"
[4] "biplot"         "boxplot"      "boxplot.default"
[7] "boxplot.matrix" "boxplot.stats" "cdplot"
[10] "coplot"         ".__C__recordedplot" "fourfoldplot"
[13] "interaction.plot" "lag.plot"     "matplot"
[16] "monthplot"      "mosaicplot"   "plot"
[19] "plot.default"   "plot.density" "plot.design"
[22] "plot.ecdf"      "plot.function" "plot.lm"
[25] "plot.mlm"       "plot.new"     "plot.spec"
[28] "plot.spec.coherency" "plot.spec.phase" "plot.stepfun"
[31] "plot.ts"        "plot.TukeyHSD" "plot.window"
[34] "plot.xy"        "preplot"      "qqplot"
[37] "recordPlot"     "replayPlot"   "savePlot"
[40] "screeplot"      "spineplot"    "sunflowerplot"
[43] "termplot"      "ts.plot"
```

- Further Reading:
 - For a useful list of “standard” packages in R, see <http://cran.fhcrc.org/doc/contrib/refcard.pdf>
 - Browse through the index of help pages in specific packages, produced by


```
> library(help = base)
```

```
> library(help = graphics)
```

 etc., and read the topics that seem interesting.

Next...

- Importing data
- Packages
- Session management and serialization
- Expressions and Objects
- Functions
- Special values
- Vectors, Matrices, Arrays, Factors, Lists, Data Frames
- Indexing,
- Logical comparisons
- Modifying Objects
- Sorting